

SHREYASH UMESH KOLHE

Software Engineer — Backend & Applied AI

Pune, Maharashtra | shreyashkolhe23@gmail.com | 9049331922 | LinkedIn | Github

EDUCATION

Master of Computer Applications (MCA) <i>Savitribai Phule Pune University</i>	2024 – 2026
Bachelor of Computer Applications (BCA) <i>Sant Gadge Baba Amravati University</i>	2021 – 2024

EXPERIENCE

Backend Developer Intern <i>CareerPhi (Remote)</i>	Sep 2025 – Dec 2025
- Led system design and optimized RESTful APIs in Java and Spring Boot, achieving a 40% reduction in average response latency for 10K+ daily requests. - Reduced request payload processing overhead by 35% through query optimization and leaner service-layer design, cutting P95 response times from 800ms to 520ms. - Implemented clean, layered service architecture to improve scalability and long-term maintainability.	

TECHNICAL SKILLS

AI & Agents: LangChain, LangGraph, RAG, Semantic Search, Vector Databases, Embeddings, Prompt Guardrails & Evals
Backend & Systems: Java (Spring Boot), Node.js (NestJS), Python, REST APIs, Microservices, Kafka, Redis, WebSocket
Databases: PostgreSQL, MySQL, MongoDB
Cloud & Infra: AWS (EC2, RDS, Elastic Beanstalk, Lambda, ECS, ECR), Docker
Frontend: React.js, HTML5, CSS3
Tools & Practices: Git, Postman, CI/CD, Agile/Scrum, Design Patterns, OOP

PROJECTS

LiftNShift — Full-Stack Home Relocation Platform <i>Personal Project</i>	GitHub Live Demo
Java · Spring Boot · React 18 · MySQL · Redis · Kafka · Razorpay · JWT · Gemini AI · OpenRouteService	
- Built full-stack relocation platform with GPS-based dynamic pricing, slot booking, and JWT auth (User/Admin); integrated Razorpay with HMAC-SHA256 verification. - Architected event-driven post-payment flow using Kafka (3 consumers) and Gemini AI for truck recommendations, decoupling booking confirmation from downstream tasks. - Eliminated N+1 queries (51→1 call/request), cut DB load ~60% via Redis caching, and prevented race conditions in slot booking with DB-level constraints + <code>@Retryable</code>.	
AI Voice Agent — Live in Production <i>Personal Project</i>	Live Demo
Python · LangGraph · RAG · Deepgram STT · TTS · OpenAI API · Claude API · Node.js (NestJS) · WebSocket · PostgreSQL · Docker · AWS	
- Built and shipped an end-to-end AI voice agent to production; LangGraph orchestrates conversation state, tool use, and context across concurrent sessions. RAG pipeline grounded strictly in verified domain data to meet real accuracy constraints, with real-time STT/TTS over WebSocket. - Implemented RBAC, full audit logging, prompt guardrails, and a red-team eval pipeline — zero hallucination rate in production. - Built the eval harness from scratch: adversarial test cases, LLM-as-judge pipelines, and regression detection before every deployment.	